

# Design Documentation - MonteCarloSimulator Library

## **Design Documentation - MonteCarloSimulator Library**

### **Table of Contents**

- 1. High-Level Overview
  - 1.1 Design Philosophy
  - 1.2 Architecture Summary
  - 1.3 Key Design Goals
- 2. Core Components
  - 2.1 SimulationEngine
  - 2.2 Execution Policies
  - 2.3 Aggregators
  - 2.4 Transforms
  - 2.5 Random Number Generation
- 3. Concept-Driven Design
  - 3.1 SimulationModel Concept
  - 3.2 ResultAggregator Concept
  - 3.3 Transform Concept
  - 3.4 RngFactory Concept
- 4. Extensibility Points
  - 4.1 Custom Models
  - 4.2 Custom Execution Policies
  - 4.3 Custom Aggregators
  - 4.4 Custom Transforms
  - 4.5 Custom RNG Factories
- 5. Performance Considerations
  - 5.1 Parallel Execution Design
  - 5.2 Memory Layout and Cache Efficiency
- 6. Build System Architecture
  - 6.1 CMake Configuration
  - 6.2 Header-Only Design
  - 6.3 Optional Features

- 7. Error Handling and Safety
  - 8. Future Extensions
  - 9. Design Patterns
  - 10. Usage Workflows
- 

## 1. High-Level Overview

### 1.1 Design Philosophy

MonteCarloSimulator is designed around the following core principles:

1. **Zero-Cost Abstractions**: Leverage C++20 features (concepts, templates) to provide compile-time polymorphism with no runtime overhead
2. **Composability**: Each component (model, execution policy, aggregator, transform) is independently configurable and composable
3. **Type Safety**: Use concepts to enforce compile-time constraints on user-provided types
4. **Header-Only**: Simplify integration and enable maximum compiler optimization opportunities
5. **Progressive Disclosure**: Simple use cases require minimal code; advanced features are available when needed

### 1.2 Architecture Summary

The library follows a layered architecture:

The architecture is based on **template parameterization** where the **SimulationEngine** acts as a coordinator between four orthogonal concerns:

- **Model**: Defines the simulation logic (what to compute per trial)
- **Execution Policy**: Defines how trials are executed (sequential, parallel, GPU)
- **Aggregator**: Defines how results are collected and summarized
- **Transform**: Defines post-processing of individual trial results

### 1.3 Key Design Goals

Goal	Design Decision	Rationale
<b>Performance</b>	Header-only + templates	Enables inlining and optimization across boundaries
<b>Flexibility</b>	Policy-based design	Users can customize each dimension independently
<b>Correctness</b>	Concept constraints	Compile-time validation of user types

Goal	Design Decision	Rationale
<b>Reproducibility</b>	Explicit seed management	Deterministic results for testing and debugging
<b>Scalability</b>	Multiple execution policies	From single-threaded to GPU acceleration

---

## 2. Core Components

### 2.1 SimulationEngine

The `SimulationEngine` is the central orchestrator that coordinates all aspects of a Monte Carlo simulation.

#### Type Signature

```
template<
    typename Model,
    typename Aggregator = WelfordAggregator<>,
    typename ExecutionPolicy = execution::Sequential,
    typename Transform = transform::Identity,
    typename RngFactory = DefaultRngFactory
>
class SimulationEngine;
```

#### Responsibilities

1. **Lifecycle Management:** Owns instances of model, policy, transform, and RNG factory
2. **Timing:** Measures execution time for performance analysis
3. **Result Construction:** Packages aggregated statistics into a `Result` object
4. **Seed Management:** Maintains base seed and passes it to execution policies

#### Key Methods

```
Result run(std::uint64_t iterations) const;
Result simulate(std::uint64_t iterations, std::uint64_t seed) const;
std::uint64_t seed() const noexcept;
void set_seed(std::uint64_t seed) noexcept;
```

#### Design Rationale




Figure 1: Architecture Layers

- **Template Parameters:** All components are template parameters rather than runtime polymorphism (virtual functions) to enable zero-cost abstraction
- **Const Correctness:** `run()` is const to allow reuse of the same engine instance
- **Seed Override:** `simulate()` allows per-run seed override for convenience while maintaining the base seed

## 2.2 Execution Policies

Execution policies determine **how** trials are executed. They implement a common interface but use static (compile-time) polymorphism.

### Interface Contract

```
template<typename Model, typename Aggregator, typename RngFactory>
void run(Model&& model, Aggregator& agg, size_t iterations,
         uint64_t seed, RngFactory rng_factory) const;
```

### Available Policies

#### 2.2.1 Sequential Execution Location: `include/montecarlo/execution/sequential.hpp`

**Characteristics:** - Single-threaded execution - Simplest implementation - Deterministic ordering - Lowest overhead

**Use Cases:** - Small problem sizes (< 100k trials) - Debugging and testing - When reproducibility is critical - Models with very fast trial execution

#### Implementation Notes:

```
class Sequential {
    // Single RNG instance, sequential iteration
    // No synchronization overhead
};
```

#### 2.2.2 Parallel Execution Location: `include/montecarlo/execution/parallel.hpp`

**Characteristics:** - Multi-threaded using `std::thread` - Work partitioning across cores - Independent RNG streams per thread - Aggregator merging strategy

#### Design Details:

1. **Thread Pool:** Creates `num_threads` worker threads (defaults to `hardware_concurrency()`)
2. **Work Distribution:** Evenly distributes iterations with remainder handling




Figure 2: SimulationEngine Interaction

3. **RNG Independence:** Each thread gets unique seed: `base_seed + thread_id`
4. **Aggregation:** Local aggregators per thread, merged at the end

Merge Strategy:

```
// Prefer native merge if available (e.g., WelfordAggregator)
if constexpr (requires { agg.merge(local_agg); }) {
    agg.merge(local_agg);
} else {
    // Fallback: replay individual results
    for (size_t i = 0; i < local_agg.count(); ++i) {
        agg.add(local_agg.result());
    }
}
```

**Use Cases:** - Large problem sizes ( $> 100k$  trials) - CPU-bound models - Multi-core systems - When speedup outweighs thread overhead




Figure 3: Parallel Execution Flow

**2.2.3 GPU Execution (Experimental)** **Location:** `include/montecarlo/execution/gpu.hpp`

**Status:** Infrastructure only, implementation pending

**Design Considerations:** - CUDA kernel execution - Device memory management - RNG on GPU (cuRAND) - Reduction operations for aggregation

**Future Implementation Notes:** - Will require CUDA-compatible model functions - Likely use thrust library for reductions - Need careful memory transfer optimization

### 2.3 Aggregators

Aggregators collect trial results and compute summary statistics. They must satisfy the `ResultAggregator` concept.

**2.3.1 WelfordAggregator Algorithm:** Welford's online variance algorithm

**Advantages:** - **Numerically Stable:** Avoids catastrophic cancellation - **Single-Pass:** Computes mean and variance in one pass - **Memory Efficient:**  $O(1)$  space complexity - **Parallelizable:** Supports Chan's parallel variance algorithm for merging

**Mathematical Foundation:**

For running mean and variance:

$$\delta = x_n - \mu_{n-1} \quad (1)$$

$$\mu_n = \mu_{n-1} + \frac{\delta}{n} \quad (2)$$

$$M_{2,n} = M_{2,n-1} + \delta \cdot (x_n - \mu_n) \quad (3)$$

$$\sigma^2 = \frac{M_2}{n-1} \quad (4)$$

**Merge Operation** (Chan's Algorithm):

```
void merge(const WelfordAggregator& other) {
    double total = count_ + other.count_;
    double delta = other.mean_ - mean_;
    mean_ += delta * (other.count_ / total);
    m2_ += other.m2_ + delta * delta * (count_ * other.count_ / total);
    count_ += other.count_;
}
```

**Use Cases:** - Default choice for most simulations - When variance/standard error is needed - Parallel execution (supports efficient merging)

**2.3.2 HistogramAggregator Purpose:** Collect distribution information

**Configuration:**

```
HistogramAggregator(size_t bins, double min, double max);
```

**Characteristics:** - Fixed-width binning -  $O(\text{bins})$  space complexity - Useful for distribution visualization - Does not compute variance

**Use Cases:** - Analyzing result distributions - Quality control checks - Detecting outliers or anomalies

**Design Trade-off:** Memory vs. information richness

## 2.4 Transforms

Transforms apply post-processing to individual trial results before aggregation.

**Design Rationale Why Transform Before Aggregation?** 1. **Avoid Double Work:** Transform once per trial, not once per final result 2. **Distribution Analysis:** Aggregator sees transformed values 3. **Composability:** Different transforms for different analyses

### Available Transforms

Transform	Formula	Use Case
Identity	$f(x) = x$	Default (no transformation)
Square	$f(x) = x^2$	Variance estimation
Abs	$f(x) =  x $	Absolute errors
Log	$f(x) = \ln(x + \text{offset})$	Log-scale analysis
Exp	$f(x) = e^x$	Exponential scaling
Indicator	$f(x) = 1 \text{ if condition else } 0$	Probability estimation
Clamp	$f(x) = \text{clamp}(x, \text{min}, \text{max})$	Range limiting
LinearScale	$f(x) = ax + b$	Rescaling outputs
Power	$f(x) = x^p$	Power transformations
Sigmoid	$f(x) = 1/(1 + e^{-x})$	Normalization

### Implementation Pattern

```
struct SomeTransform {
    double operator()(double x) const noexcept {
        // Transformation logic
        return transformed_value;
    }
};
```

**Example Use Case: Indicator Function** Estimating probability  $P(X > \text{threshold})$ :

```
auto engine = make_sequential_engine(
    some_model,
    123456789ULL,
    transform::Indicator{threshold, true} // Returns 1 if > threshold
);
```

```
auto result = engine.run(1000000);
// result.estimate is roughly P(X > threshold)
```

## 2.5 Random Number Generation

### Design Goals

1. **Reproducibility:** Same seed → same sequence
2. **Independence:** Parallel streams must be uncorrelated
3. **Flexibility:** Users can inject custom RNGs
4. **Quality:** Use well-tested algorithms

### RNG Factory Pattern Interface:

```
template<typename F>
concept RngFactory = requires(F f, std::uint64_t s) {
    { f(s) };
} && std::uniform_random_bit_generator<
    std::decay_t<decltype(std::declval<F>()> Ou)>
>;
```

### Default Implementation:

```
struct DefaultRngFactory {
    std::mt19937_64 operator()(std::uint64_t seed) const {
        return make_rng(seed);
    }
};
```

**Seed Sequence Strategy** To ensure independent streams in parallel execution:

```
inline std::mt19937_64 make_rng(std::uint64_t seed,
                                  std::uint64_t stream_id = 0) {
    std::seed_seq seq{
        static_cast<unsigned>(seed),
        static_cast<unsigned>(seed >> 32),
        static_cast<unsigned>(stream_id),
        static_cast<unsigned>(stream_id >> 32)
    };
    return std::mt19937_64(seq);
}
```

**Design Rationale:** - Uses `std::seed_seq` to mix seed and `stream_id` - Ensures different threads get different sequences - Avoids correlation issues from simple seed arithmetic

---

### 3. Concept-Driven Design

C++20 concepts provide **compile-time contracts** that:

1. Improve error messages
2. Enable SFINAE-free overloading
3. Document requirements
4. Prevent misuse

#### 3.1 SimulationModel Concept

```
template<typename Model, typename RNG>
concept SimulationModel = TrialModel<Model, RNG> || CallableModel<Model, RNG>;

template<typename Model, typename RNG>
concept TrialModel = requires(Model m, RNG& rng) {
    { m.trial(rng) } -> std::convertible_to<double>;
};

template<typename Model, typename RNG>
concept CallableModel = requires(Model m, RNG& rng) {
    { m(rng) } -> std::convertible_to<double>;
};
```

**Design Decision:** Support both `.trial()` method and `operator()` for flexibility.

**Benefits:**

- Users can choose their preferred interface
- Lambdas work naturally with `operator()`
- Class-based models can use `.trial()` for clarity

**Detection at Runtime:**

```
if constexpr (requires { model.trial(rng); }) {
    result = model.trial(rng);
} else {
    result = model(rng);
}
```

#### 3.2 ResultAggregator Concept

```
template<typename Aggregator, typename T>
concept ResultAggregator = requires(Aggregator agg, T value) {
    { agg.add(value) } -> std::same_as<void>;
    { agg.result() } -> std::convertible_to<double>;
    { agg.reset() } -> std::same_as<void>;
};
```

**Required Methods:**

- `add(value)`: Accumulate a new result
- `result()`: Return current estimate
- `reset()`: Clear state for reuse

**Optional Methods** (detected via SFINAE):

- `variance()`: Return sample variance
- `std_error()`: Return standard error
- `merge(other)`: Merge another aggregator
- `count()`: Return number of samples

### 3.3 Transform Concept

```
template<typename T>
concept Transform = requires(T transform, double value) {
    { transform(value) } -> std::convertible_to<double>;
};
```

**Simplicity:** Just a callable taking `double` → `double`

**Stateless vs. Stateful:** - Stateless (e.g., `Identity`, `Square`): Simple functors - Stateful (e.g., `LinearScale`, `Clamp`): Store parameters

### 3.4 RngFactory Concept

```
template<typename F>
concept RngFactory = requires(F f, std::uint64_t s) {
    { f(s) };
} && std::uniform_random_bit_generator<
    std::decay_t<decltype(std::declval<F>()(0u))>
>;
```

**Requirements:** 1. Callable with `uint64_t` seed 2. Returns a type satisfying `std::uniform_random_bit_generator`

**Use Case:** Enables custom RNGs for specialized needs (e.g., testing, SIMD, hardware RNGs)

---

## 4. Extensibility Points

### 4.1 Custom Models

**Example:** Multi-dimensional integration

```
class MultiDimIntegrationModel {
    std::function<double(const std::vector<double>&)> func_;
    size_t dimensions_;

public:
    MultiDimIntegrationModel(auto f, size_t dim)
        : func_(f), dimensions_(dim) {}

    template<typename RNG>
    double trial(RNG& rng) {
        std::uniform_real_distribution<> dist(0.0, 1.0);
        std::vector<double> point(dimensions_);
        for (auto& coord : point) {
            coord = dist(rng);
        }
    }
}
```

```

        return func_(point);
    }
};
```

## 4.2 Custom Execution Policies

Requirements:

```
template<typename Model, typename Aggregator, typename RngFactory>
void run(Model&& model, Aggregator& agg, size_t iterations,
         uint64_t seed, RngFactory rng_factory) const;
```

**Example:** Batched execution with checkpointing

```
class CheckpointedExecution {
    std::filesystem::path checkpoint_path_;
    size_t batch_size_;

public:
    template<typename Model, typename Aggregator, typename RngFactory>
    void run(Model&& model, Aggregator& agg, size_t iterations,
             uint64_t seed, RngFactory rng_factory) const {
        for (size_t i = 0; i < iterations; i += batch_size_) {
            // Execute batch
            // Save checkpoint
        }
    }
};
```

## 4.3 Custom Aggregators

**Example:** Quantile estimation

```
class QuantileAggregator {
    std::vector<double> samples_;

public:
    void add(double value) {
        samples_.push_back(value);
    }

    double result() const {
        return median();
    }

    double quantile(double q) const {
        auto sorted = samples_;
        std::sort(sorted.begin(), sorted.end());
```

```

        size_t idx = static_cast<size_t>(q * sorted.size());
        return sorted[idx];
    }

    void reset() {
        samples_.clear();
    }
};


```

**Note:** Memory-intensive for large sample counts. Consider streaming algorithms (P<sup>2</sup> algorithm) for production use.

#### 4.4 Custom Transforms

**Example:** Discounting future values

```

struct DiscountTransform {
    double rate;
    double time;

    double operator()(double future_value) const {
        return future_value * std::exp(-rate * time);
    }
};

```

#### 4.5 Custom RNG Factories

**Example:** Testing with deterministic sequence

```

class ConstantRngFactory {
    double value_;

public:
    explicit ConstantRngFactory(double val) : value_(val) {}

    auto operator()(std::uint64_t) const {
        return [val = value_]() mutable -> std::uint64_t {
            return static_cast<uint64_t>(val * 1e9);
        };
    }
};

```

---

### 5. Performance Considerations

#### 5.1 Parallel Execution Design

Work Distribution Strategy Even Distribution:

```

size_t iters_per_thread = iterations / num_threads_;
size_t remaining = iterations % num_threads_;

for (size_t t = 0; t < num_threads_; ++t) {
    size_t thread_iters = iters_per_thread + (t < remaining ? 1 : 0);
    // Assign to thread t
}

```

**Design Choice:** Distribute remainder to first threads rather than assigning to last thread (better load balance when remainder is large).

**Synchronization Overhead Zero Synchronization During Execution:** -  
Each thread has independent RNG - Each thread has independent aggregator -  
No shared state during trial execution

**Single Synchronization Point:** - Thread join at end - Aggregator merge is sequential

**Performance Model:**  $T_{parallel} = \frac{T_{sequential}}{N_{threads}} + T_{merge} + T_{overhead}$

Where: -  $T_{merge}$  is  $O(\text{num\_threads})$  for Welford merge -  $T_{overhead}$  is thread creation cost (amortized over many iterations)

**Speedup Analysis Amdahl's Law Applied:**

Speedup =  $1 / (f + (1-f)/N)$

Where: -  $f$  = fraction of serial work (merge + overhead) -  $N$  = number of threads

For typical Monte Carlo: -  $f$  is approximately 0.001 (merge is negligible) -  
Theoretical speedup is approximately  $N$  for large iteration counts

## 5.2 Memory Layout and Cache Efficiency

### Engine Memory Layout

SimulationEngine object:  
 |- model\_ [sizeof(Model)]  
 |- policy\_ [sizeof(ExecutionPolicy)]  
 |- transform\_ [sizeof(Transform)]  
 |- rng\_factory\_ [sizeof(RngFactory)]  
 `-- base\_seed\_ [8 bytes]

**Small Object Optimization:** - Stateless types (Identity transform, Default-RngFactory) = empty base optimization - Total size typically 16-32 bytes for default configuration

**Aggregator Memory Access Pattern Sequential Execution:** - Cache-friendly: single aggregator, sequential updates - No cache thrashing

- Parallel Execution:**
- Each thread updates local aggregator (no false sharing)
  - Aggregators are cache-line aligned (implicit via separate objects)

## 6. Build System Architecture

### 6.1 CMake Configuration

#### Option-Based Feature Control

```
option(MCLIB_BUILD_EXAMPLES "Build example applications" ON)
option(MCLIB_BUILD_TESTS "Build unit tests" ON)
option(MCLIB_ENABLE_PARALLEL "Enable parallel CPU execution" ON)
option(MCLIB_ENABLE_GPU "Enable GPU acceleration (CUDA)" OFF)
```

**Design Philosophy:** Features are opt-in with sensible defaults

#### Compile Definitions

CMake Option	C++ Macro	Effect
MCLIB_ENABLE_PARALLEL=ON	MCLIB_PARALLEL_ENABLED	Includes parallel.hpp, enables make_parallel_engine()
MCLIB_ENABLE_GPU=ON	MCLIB_GPU_ENABLED	Includes gpu.hpp, enables CUDA support

#### Target Configuration

```
add_library(montecarlo INTERFACE) # Header-only library

target_compile_features(montecarlo INTERFACE cxx_std_20)

target_include_directories(montecarlo INTERFACE
    ${BUILD_INTERFACE:${CMAKE_CURRENT_SOURCE_DIR}/include}
    ${INSTALL_INTERFACE:include}
)
```

**Generator Expressions:** Separate build-time and install-time include paths

### 6.2 Header-Only Design

#### Benefits

1. **No ABI Issues:** Template instantiation at usage site
2. **Maximum Optimization:** Compiler sees full implementation
3. **Easy Integration:** Just add include directory
4. **No Linking:** Simplifies build system

### Trade-offs

1. **Compile Time:** Template instantiation cost
2. **Code Bloat:** Multiple instantiations possible
3. **Debug Symbols:** Can be large for complex templates

**Mitigation:** - Extern template declarations (future work) - Precompiled headers for common instantiations

### 6.3 Optional Features

Dependency Management Parallel Support:

```
if(MCLIB_ENABLE_PARALLEL)
    find_package(Threads REQUIRED)
    target_link_libraries(montecarlo INTERFACE Threads::Threads)
endif()
```

GPU Support (Placeholder):

```
if(MCLIB_ENABLE_GPU)
    enable_language(CUDA)
    find_package(CUDAToolkit REQUIRED)
    target_link_libraries(montecarlo INTERFACE CUDA::cudart)
endif()
```

Conditional Compilation In Headers:

```
#ifdef MCLIB_PARALLEL_ENABLED
#include "execution/parallel.hpp"
#endif
```

In Factory Functions:

```
#ifdef MCLIB_PARALLEL_ENABLED
template<typename Model>
auto make_parallel_engine(Model model, size_t threads = 0) {
    // Implementation
}
#endif
```

---

## 7. Error Handling and Safety

Compile-Time Errors

**Concept Violations:** Clear error messages when requirements aren't met

```
// Error if Model doesn't satisfy SimulationModel
static_assert(SimulationModel<Model, std::mt19937_64>);
```

## Runtime Errors

**Current Approach:** Minimal runtime checks (rely on correct usage)

## Safety Guarantees

**No Undefined Behavior** (assuming valid models): - No raw pointer arithmetic  
 - No manual memory management - RAII for all resources - Thread-safe by design (no shared mutable state during execution)

---

## 8. Future Extensions

### Planned Features

1. **Adaptive Sampling:** Stop when error target is reached

```
Result run_until_error(double target_error, size_t max_iterations);
```

2. **GPU Implementation:** Full CUDA kernel support

3. **Distributed Computing:** MPI-based execution policy

### API Stability

**Current Status:** Release (1.0)

**Compatibility Promise** (post-1.0): - Core API (SimulationEngine, concepts) will remain stable - Execution policies may be extended but won't break existing code - New aggregators/transforms can be added without breaking changes

---

## 9. Design Patterns

### Strategy Pattern

**Application:** Execution policies - Different algorithms (sequential, parallel) - Same interface - Runtime selection via template parameter

### Template Method Pattern

**Application:** SimulationEngine::run() - Defines skeleton (timing, wrapping, result construction) - Delegates trial execution to policy

### Factory Pattern

**Application:** Engine creation helpers

```
make_sequential_engine(model);
make_parallel_engine(model, threads);
```

- Simplifies complex template instantiation
- Provides sensible defaults

### Policy-Based Design

**Application:** Overall architecture - Orthogonal concerns (model, execution, aggregation, transform) - Compile-time composition - Zero runtime overhead

---

## 10. Usage Workflows

### Basic Workflow

#### Advanced Workflow

```
// 1. Define custom components
class MyModel { /* ... */ };
class MyAggregator { /* ... */ };
struct MyTransform { /* ... */ };

// 2. Compose engine
auto engine = SimulationEngine<
    MyModel,
    MyAggregator,
    execution::Parallel,
    MyTransform
>(
    my_model,
    execution::Parallel{8},   // 8 threads
    MyTransform{params},
    DefaultRngFactory{},
    seed
);

// 3. Run with custom configuration
auto result = engine.run(1'000'000);

// 4. Extract detailed statistics
double ci_lower = ci_95(result).lower;
double ci_upper = ci_95(result).upper;
```

### Testing Workflow

```
// Use deterministic RNG for tests
struct TestRngFactory {
    std::mt19937_64 operator()(uint64_t seed) const {
        return std::mt19937_64{seed};  // Fixed seed
```




Figure 4: Basic Workflow

```

    }
};

auto engine = make_engine<MyModel, Sequential, WelfordAggregator<>,
    Identity, TestRngFactory>(
    model, Sequential{}, fixed_seed, TestRngFactory{}, Identity{}
);

auto result = engine.run(1000);
ASSERT_NEAR(result.estimate, expected, tolerance);

```

---

## Conclusion

The MonteCarloSimulator library exemplifies modern C++ design principles:

- **Type Safety:** Concepts enforce correctness at compile time
- **Performance:** Zero-cost abstractions via templates and inlining
- **Flexibility:** Policy-based design allows customization of all components
- **Simplicity:** Header-only, minimal dependencies, clean API
- **Extensibility:** Clear extension points for custom behavior

The architecture supports a wide range of Monte Carlo applications while maintaining a simple mental model: define a model, choose an execution policy, run trials, and collect results. Advanced users can customize any aspect of the pipeline without sacrificing performance or type safety.